

```
In [100]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
```

```
In [101]: data = pd.read_csv(r'C:\Users\modid\Desktop\NBA_Rookie_Predictions.csv')
data.shape
#importing the dataset in and checking shape
```

```
Out[101]: (1340, 21)
```

```
In [102]: data_new = data.dropna()
data_new.shape
#dropping the rows with missing values
```

```
Out[102]: (1329, 21)
```

```
In [103]: X = data_new.drop(['name', 'target_5yrs'], axis=1)
y = data_new['target_5yrs']
#splitting the predictors and target
```

```
In [104]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
#splitting the data into train and test models
```

```
In [105]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
log_reg = LogisticRegression(max_iter = 10000)
log_reg.fit(X_train, y_train) #fit to training data
y_pred_log = log_reg.predict(X_test) #predict on test data
log_cm = confusion_matrix(y_test, y_pred_log)
log_cm
```

```
Out[105]: array([[ 84,  69],
 [ 44, 202]], dtype=int64)
```

```
In [106]: error_rate = 1 - accuracy_score(y_test, y_pred_log)
error_rate
#obtaining accuracy of model
```

```
Out[106]: 0.28320802005012535
```

```
In [107]: coefficients = log_reg.coef_[0]
coefficients
#obtaining coefficients
```

```
Out[107]: array([ 0.03608031, -0.04147289,  0.14324291, -0.72781581,  0.24059492,
 0.05091288,  1.44149439, -0.56181202,  0.00495857,  0.06826616,
-0.05954725,  0.01208377,  0.58801621, -0.27307952,  0.11206921,
 0.18578501, -0.06734523,  0.75187748,  0.06231129])
```

```
In [108]: log_rep = classification_report(y_test, y_pred_log)
log_rep
#goodness of fit
```

```
Out[108]: '          precision    recall  f1-score   support\n\n
153\n          1         0.75         0.82         0.78         246\n\n
0.72         399\n  macro avg         0.70         0.69         0.69         399\nweighted avg         0.71
```

```
In [109]: std = np.std(X_train, axis=0)
sample_size = X_train.shape[0]
se = std/np.sqrt(sample_size)
se
#standard error of each of the coefficients
```

```
Out[109]: gp          0.570715
min          0.272810
pts          0.143162
fgm          0.055426
fga          0.119436
fg           0.202989
3p_made      0.012579
3pa          0.035041
3p           0.516587
ftm          0.031636
fta          0.042296
ft           0.337288
oreb         0.024789
dreb         0.043053
reb          0.065175
ast          0.046810
stl          0.013413
blk          0.014200
tov          0.024069
dtype: float64
```

```
In [110]: from sklearn.model_selection import cross_val_score
cv_acc = cross_val_score(log_reg, X_train, y_train, cv=5, scoring='accuracy')
cv_acc
#running cross validation on the model, using the X and y training data
#doing k fold cross validation, using 5 folds
```

```
Out[110]: array([0.67204301, 0.67741935, 0.75268817, 0.70430108, 0.73655914])
```

```
In [111]: average_acc = np.mean(cv_acc)
average_acc
#taking the average of all of folds to obtain how well the model performs
```

```
Out[111]: 0.7086021505376344
```

```
In [ ]:
```

```
In [ ]:
```

```
In [112]: ###LDA Analysis
```

```
In [113]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
lda_model = LinearDiscriminantAnalysis()
lda_model.fit(X_train,y_train)
y_pred_lda = lda_model.predict(X_test)
lda_cm = confusion_matrix(y_test, y_pred_lda)
lda_cm
```

```
Out[113]: array([[ 85,  68],
 [ 43, 203]], dtype=int64)
```

```
In [114]: error_rate_lda = 1 - accuracy_score(y_test, y_pred_lda)
error_rate_lda
```

```
Out[114]: 0.27819548872180455
```

```
In [115]: coefficients_lda = lda_model.coef_[0]
coefficients_lda
```

```
Out[115]: array([ 4.12045874e-02, -3.26601457e-02,  1.13678743e-01, -1.45808228e+00,
        6.09499561e-01,  8.21856565e-02,  4.01883026e+00, -1.55291668e+00,
        2.88307629e-03, -2.09158433e-01,  1.22587172e-01,  1.70911894e-02,
        1.25150723e+00,  4.23120527e-01, -6.06899357e-01,  1.64622230e-01,
       -1.18019383e-02,  6.02904733e-01,  4.58506190e-02])
```

```
In [116]: lda_rep = classification_report(y_test, y_pred_lda)
lda_rep
```

```
Out[116]: '              precision    recall  f1-score   support\n\n         0               0.66       0.56       0.60\n153\n         1               0.75       0.83       0.79       246\n0.72       399\n         macro avg       0.71       0.69       0.70       399\n0.72       0.72       399\n'
```

```
In [117]: cv_acc_lda = cross_val_score(lda_model, X_train, y_train, cv=5, scoring='accuracy')
cv_acc_lda
```

```
Out[117]: array([0.66129032, 0.66129032, 0.75268817, 0.6827957 , 0.73655914])
```

```
In [118]: average_acc_lda = np.mean(cv_acc_lda)
average_acc_lda
```

```
Out[118]: 0.6989247311827956
```

```
In [119]: from sklearn.utils import resample
num_bootstraps = 1000
coef_list = []
```

```
In [158]: for i in range(num_bootstraps):
    X_sample, y_sample = resample(X, y)
    lda = LinearDiscriminantAnalysis()
    lda.fit(X_sample, y_sample)
    coef_list.append(lda_model.coef_[0])
```

```
coef_array = np.array(coef_list)
coef_array
```

```
Out[158]: array([[ 0.03464551, -0.02527599, -0.52279653, ..., -0.30477843,
        0.41237472,  0.11052142],
       [ 0.04698796, -0.0896188 ,  0.28757558, ...,  0.27664974,
        0.34485008, -0.30430118],
       [ 0.04301429, -0.00233044,  0.11774165, ..., -0.39032792,
       -0.16847388, -0.15700336],
       ...,
       [ 0.04120459, -0.03266015,  0.11367874, ..., -0.01180194,
        0.60290473,  0.04585062],
       [ 0.04120459, -0.03266015,  0.11367874, ..., -0.01180194,
        0.60290473,  0.04585062],
       [ 0.04120459, -0.03266015,  0.11367874, ..., -0.01180194,
        0.60290473,  0.04585062]])
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

In [28]: *###QDA(Used to be K Nearest Neighbors, but issue with new Sklearn Library, KNN results on my document)*

```
In [121]: from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
qda_model = QuadraticDiscriminantAnalysis()
qda_model.fit(X_train,y_train)
y_pred_qda = qda_model.predict(X_test)
qda_cm = confusion_matrix(y_test, y_pred_qda)
qda_cm
```

Out[121]: array([[116, 37],  
[100, 146]], dtype=int64)

```
In [122]: error_rate_qda = 1 - accuracy_score(y_test, y_pred_qda)
error_rate_qda
```

Out[122]: 0.343358395989975

In [ ]:

```
In [123]: qda_rep = classification_report(y_test, y_pred_qda)
qda_rep
```

Out[123]:

		precision	recall	f1-score	support		0	0.54	0.76	0.63
153	1	0.80	0.59	0.68	246	accuracy				
0.66	399	macro avg	0.67	0.68	0.65	399	nweighted avg		0.70	
0.66	0.66	399								

```
In [124]: cv_acc_qda = cross_val_score(qda_model, X_train, y_train, cv=5, scoring='accuracy')
cv_acc_qda
```

Out[124]: array([0.62903226, 0.62903226, 0.69892473, 0.60752688, 0.6344086 ])

```
In [125]: average_acc_qda = np.mean(cv_acc_qda)
average_acc_qda
```

Out[125]: 0.6397849462365591

In [ ]:

In [126]: *###Random Forest*

```
In [127]: from sklearn.ensemble import RandomForestClassifier
rf_model = RandomForestClassifier()
rf_model.fit(X_train, y_train)
y_pred_rf = rf_model.predict(X_test)
rf_cm = confusion_matrix(y_test, y_pred_rf)
rf_cm
```

Out[127]: array([[ 73, 80],  
[ 54, 192]], dtype=int64)

```
In [128]: error_rate_rf = 1 - accuracy_score(y_test, y_pred_rf)
error_rate_rf
```

Out[128]: 0.3358395989974937

```
In [129]: rf_rep = classification_report(y_test, y_pred_rf)
rf_rep
```

```
Out[129]: '          precision    recall  f1-score   support\n\n         0          0.57    0.48    0.52\n153\n         1          0.71    0.78    0.74    246\n0.66      399\n         macro avg    0.64    0.63    0.63    399\n0.66      0.66      399\n'
```

```
In [130]: cv_acc_rf = cross_val_score(rf_model, X_train, y_train, cv=5, scoring='accuracy')
cv_acc_rf
```

```
Out[130]: array([0.67741935, 0.70430108, 0.74193548, 0.64516129, 0.70430108])
```

```
In [131]: average_acc_rf = np.mean(cv_acc_rf)
average_acc_rf
```

```
Out[131]: 0.6946236559139786
```

```
In [132]: feature_importance = pd.DataFrame({
    'Feature': X_train.columns,
    'Coefficient': rf_model.feature_importances_})

feature_importance
```

```
Out[132]:
```

	Feature	Coefficient
0	gp	0.111924
1	min	0.064362
2	pts	0.062196
3	fgm	0.062725
4	fga	0.050448
5	fg	0.082833
6	3p_made	0.022418
7	3pa	0.032614
8	3p	0.044514
9	ftm	0.050234
10	fta	0.052687
11	ft	0.061918
12	oreb	0.044187
13	dreb	0.047242
14	reb	0.051974
15	ast	0.046681
16	stl	0.032952
17	blk	0.038293
18	tov	0.039800

```
In [133]: num_bootstraps = 10
feature_importances = []
```

```
In [134]: for i in range(num_bootstraps):
          X_sample, y_sample = resample(X, y)
          rf = RandomForestClassifier()
          rf.fit(X_sample, y_sample)
          feature_importances.append(rf.feature_importances_)
```

```
In [135]: feature_array = np.array(feature_importances)
          standard_errors = np.std(feature_array, axis=0)
          standard_errors
```

```
Out[135]: array([0.01497074, 0.00657489, 0.00511298, 0.00616035, 0.00391987,
                 0.00774345, 0.00174248, 0.00199712, 0.00321969, 0.00776175,
                 0.00451264, 0.00460588, 0.00613606, 0.00413104, 0.00363989,
                 0.00176143, 0.00234078, 0.00443688, 0.00261307])
```

```
In [ ]:
```

```
In [ ]:
```

```
In [136]: ###Lasso Regression for Classification
```

```
In [ ]:
```

```
In [137]: from sklearn.linear_model import LogisticRegressionCV
          lasso = LogisticRegressionCV(cv=5, penalty='l1', solver='liblinear', max_iter = 100000)##automatically
          lasso.fit(X_train, y_train)
          y_pred_lasso = lasso.predict(X_test)
          lasso_cm = confusion_matrix(y_test, y_pred_lasso)
          lasso_cm
```

```
Out[137]: array([[ 85,  68],
                 [ 48, 198]], dtype=int64)
```

```
In [138]: error_rate_lasso = 1 - accuracy_score(y_test, y_pred_lasso)
          error_rate_lasso
```

```
Out[138]: 0.2907268170426065
```

```
In [139]: lasso_rep = classification_report(y_test, y_pred_lasso)
          lasso_rep
```

```
Out[139]: '          precision    recall  f1-score   support\n\n
153\n          1          0.74          0.80          0.77          246\n\n
0.71          399\n  macro avg          0.69          0.68          0.68          399\nweighted avg          0.70'
```

```
In [140]: coefficients_lasso = lasso.coef_[0]
coefficients_df = pd.DataFrame({
    'Feature': X_train.columns,
    'Coefficient': coefficients_lasso
})

coefficients_df
```

Out[140]:

	Feature	Coefficient
0	gp	0.036056
1	min	-0.048045
2	pts	0.000000
3	fgm	-0.555065
4	fga	0.313201
5	fg	0.044278
6	3p_made	3.508661
7	3pa	-1.265343
8	3p	0.002964
9	ftm	0.141796
10	fta	0.000000
11	ft	0.011798
12	oreb	0.709830
13	dreb	-0.159098
14	reb	0.000000
15	ast	0.212925
16	stl	0.000000
17	blk	0.799254
18	tov	0.000000

```
In [141]: cv_acc_lasso = cross_val_score(lasso, X_train, y_train, cv=5, scoring='accuracy')
cv_acc_lasso
```

Out[141]: array([0.66666667, 0.64516129, 0.75268817, 0.67741935, 0.73655914])

```
In [142]: average_acc_lasso = np.mean(cv_acc_lasso)
average_acc_lasso
```

Out[142]: 0.6956989247311828

```
In [143]: num_bootstraps = 5
bootstrapped_coefficients = []
```

```
In [144]: for i in range(num_bootstraps):
    X_sample, y_sample = resample(X_train, y_train)
    lasso = LogisticRegressionCV(cv=5, penalty='l1', solver='liblinear', max_iter=100000)
    lasso.fit(X_sample, y_sample)
    bootstrapped_coefficients.append(lasso.coef_[0])
```

```
In [145]: coefficients_array = np.array(bootstrapped_coefficients)
sd = np.std(coefficients_array, axis=0)
se = sd / np.sqrt(X_train.shape[0])
```

In [146]: standard\_errors

Out[146]: array([0.01497074, 0.00657489, 0.00511298, 0.00616035, 0.00391987,  
0.00774345, 0.00174248, 0.00199712, 0.00321969, 0.00776175,  
0.00451264, 0.00460588, 0.00613606, 0.00413104, 0.00363989,  
0.00176143, 0.00234078, 0.00443688, 0.00261307])

In [ ]:

In [ ]:

In [147]: ###PCA

In [148]: from sklearn.decomposition import PCA

In [149]: pca = PCA(n\_components=3)  
X\_train\_pca = pca.fit\_transform(X\_train)  
X\_test\_pca = pca.transform(X\_test)

In [150]: classifier = LogisticRegression()  
classifier.fit(X\_train\_pca, y\_train)

Out[150]: 

▼ LogisticRegression

  
LogisticRegression()

In [151]: y\_pred\_pca= classifier.predict(X\_test\_pca)

In [152]: pca\_cm = confusion\_matrix(y\_test, y\_pred\_pca)  
pca\_cm

Out[152]: array([[ 79, 74],  
[ 46, 200]], dtype=int64)

In [153]: error\_rate\_PCA = 1 - accuracy\_score(y\_test, y\_pred\_pca)  
error\_rate\_PCA

Out[153]: 0.3007518796992481

In [154]: pca\_rep = classification\_report(y\_test, y\_pred\_pca)  
pca\_rep

Out[154]: ' precision recall f1-score support\n\n 0 0.63 0.52 0.57  
153\n 1 0.73 0.81 0.77 246\n\n accuracy  
0.70 399\n macro avg 0.68 0.66 0.67 399\nweighted avg 0.69  
0.70 0.69 399\n'

In [155]: coefficients\_pca = classifier.coef\_[0]  
coefficients\_pca

Out[155]: array([-0.04577693, -0.02199835, 0.00417098])

In [156]: cv\_acc\_pca = cross\_val\_score(classifier, X\_train\_pca, y\_train, cv=5, scoring='accuracy')  
cv\_acc\_pca

Out[156]: array([0.69892473, 0.66666667, 0.72043011, 0.66666667, 0.7311828 ])



```
In [157]: average_acc_pca = np.mean(cv_acc_pca)
average_acc_pca
```

```
Out[157]: 0.696774193548387
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [56]:
```

```
Requirement already satisfied: scikit-learn in c:\users\modid\anaconda3\lib\site-packages (1.3.2)
Requirement already satisfied: joblib>=1.1.1 in c:\users\modid\anaconda3\lib\site-packages (from scikit-learn) (1.3.2)
Requirement already satisfied: scipy>=1.5.0 in c:\users\modid\anaconda3\lib\site-packages (from scikit-learn) (1.10.1)
Requirement already satisfied: numpy<2.0,>=1.17.3 in c:\users\modid\anaconda3\lib\site-packages (from scikit-learn) (1.24.4)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\modid\anaconda3\lib\site-packages (from scikit-learn) (2.1.0)
```

```
In [57]: ###Polynomial Regression, I was able to figure it out but did include on analysis.
```

```
In [58]: from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LogisticRegression
poly_reg = PolynomialFeatures(degree = 3)
X_train_poly = poly_reg.fit_transform(X_train)
X_test_poly = poly_reg.transform(X_test)
log_poly = LogisticRegression(solver='liblinear', max_iter=10000)
log_poly.fit(X_train_poly, y_train)
y_pred_poly = log_poly.predict(X_test_poly)
poly_cm = confusion_matrix(y_test, y_pred_poly)
poly_cm
```

```
Out[58]: array([[ 78,  81],
               [ 64, 176]], dtype=int64)
```

```
In [59]: error_rate_poly = 1 - accuracy_score(y_test, y_pred_poly)
error_rate_poly
```

```
Out[59]: 0.3634085213032582
```

```
In [60]: poly_rep = classification_report(y_test, y_pred_poly)
poly_rep
```

```
Out[60]: '          precision    recall  f1-score   support\n\n         0          0.55      0.49      159\n         1          0.68      0.73      0.71      240\n\n accuracy\n0.64      399\n macro avg      0.62      0.61      0.61      399\nweighted avg      0.63      399\n'
```

```
In [61]: coefficients_poly = log_poly.coef_[0]
coefficients_poly
```

```
Out[61]: array([-3.93924441e-05, -4.27216482e-04, -3.53518923e-04, ...,
                1.68082157e-04,  1.29783772e-04, -1.72661155e-04])
```

```
In [62]: cv_acc_poly = cross_val_score(log_poly, X_train_poly, y_train, cv=5, scoring='accuracy')
cv_acc_poly
```

```
Out[62]: array([0.60752688, 0.61827957, 0.61827957, 0.69354839, 0.6344086 ])
```

```
In [63]: average_acc_lasso = np.mean(cv_acc_lasso)
         average_acc_lasso
```

```
Out[63]: 0.7129032258064518
```

```
In [ ]:
```

```
In [ ]:
```